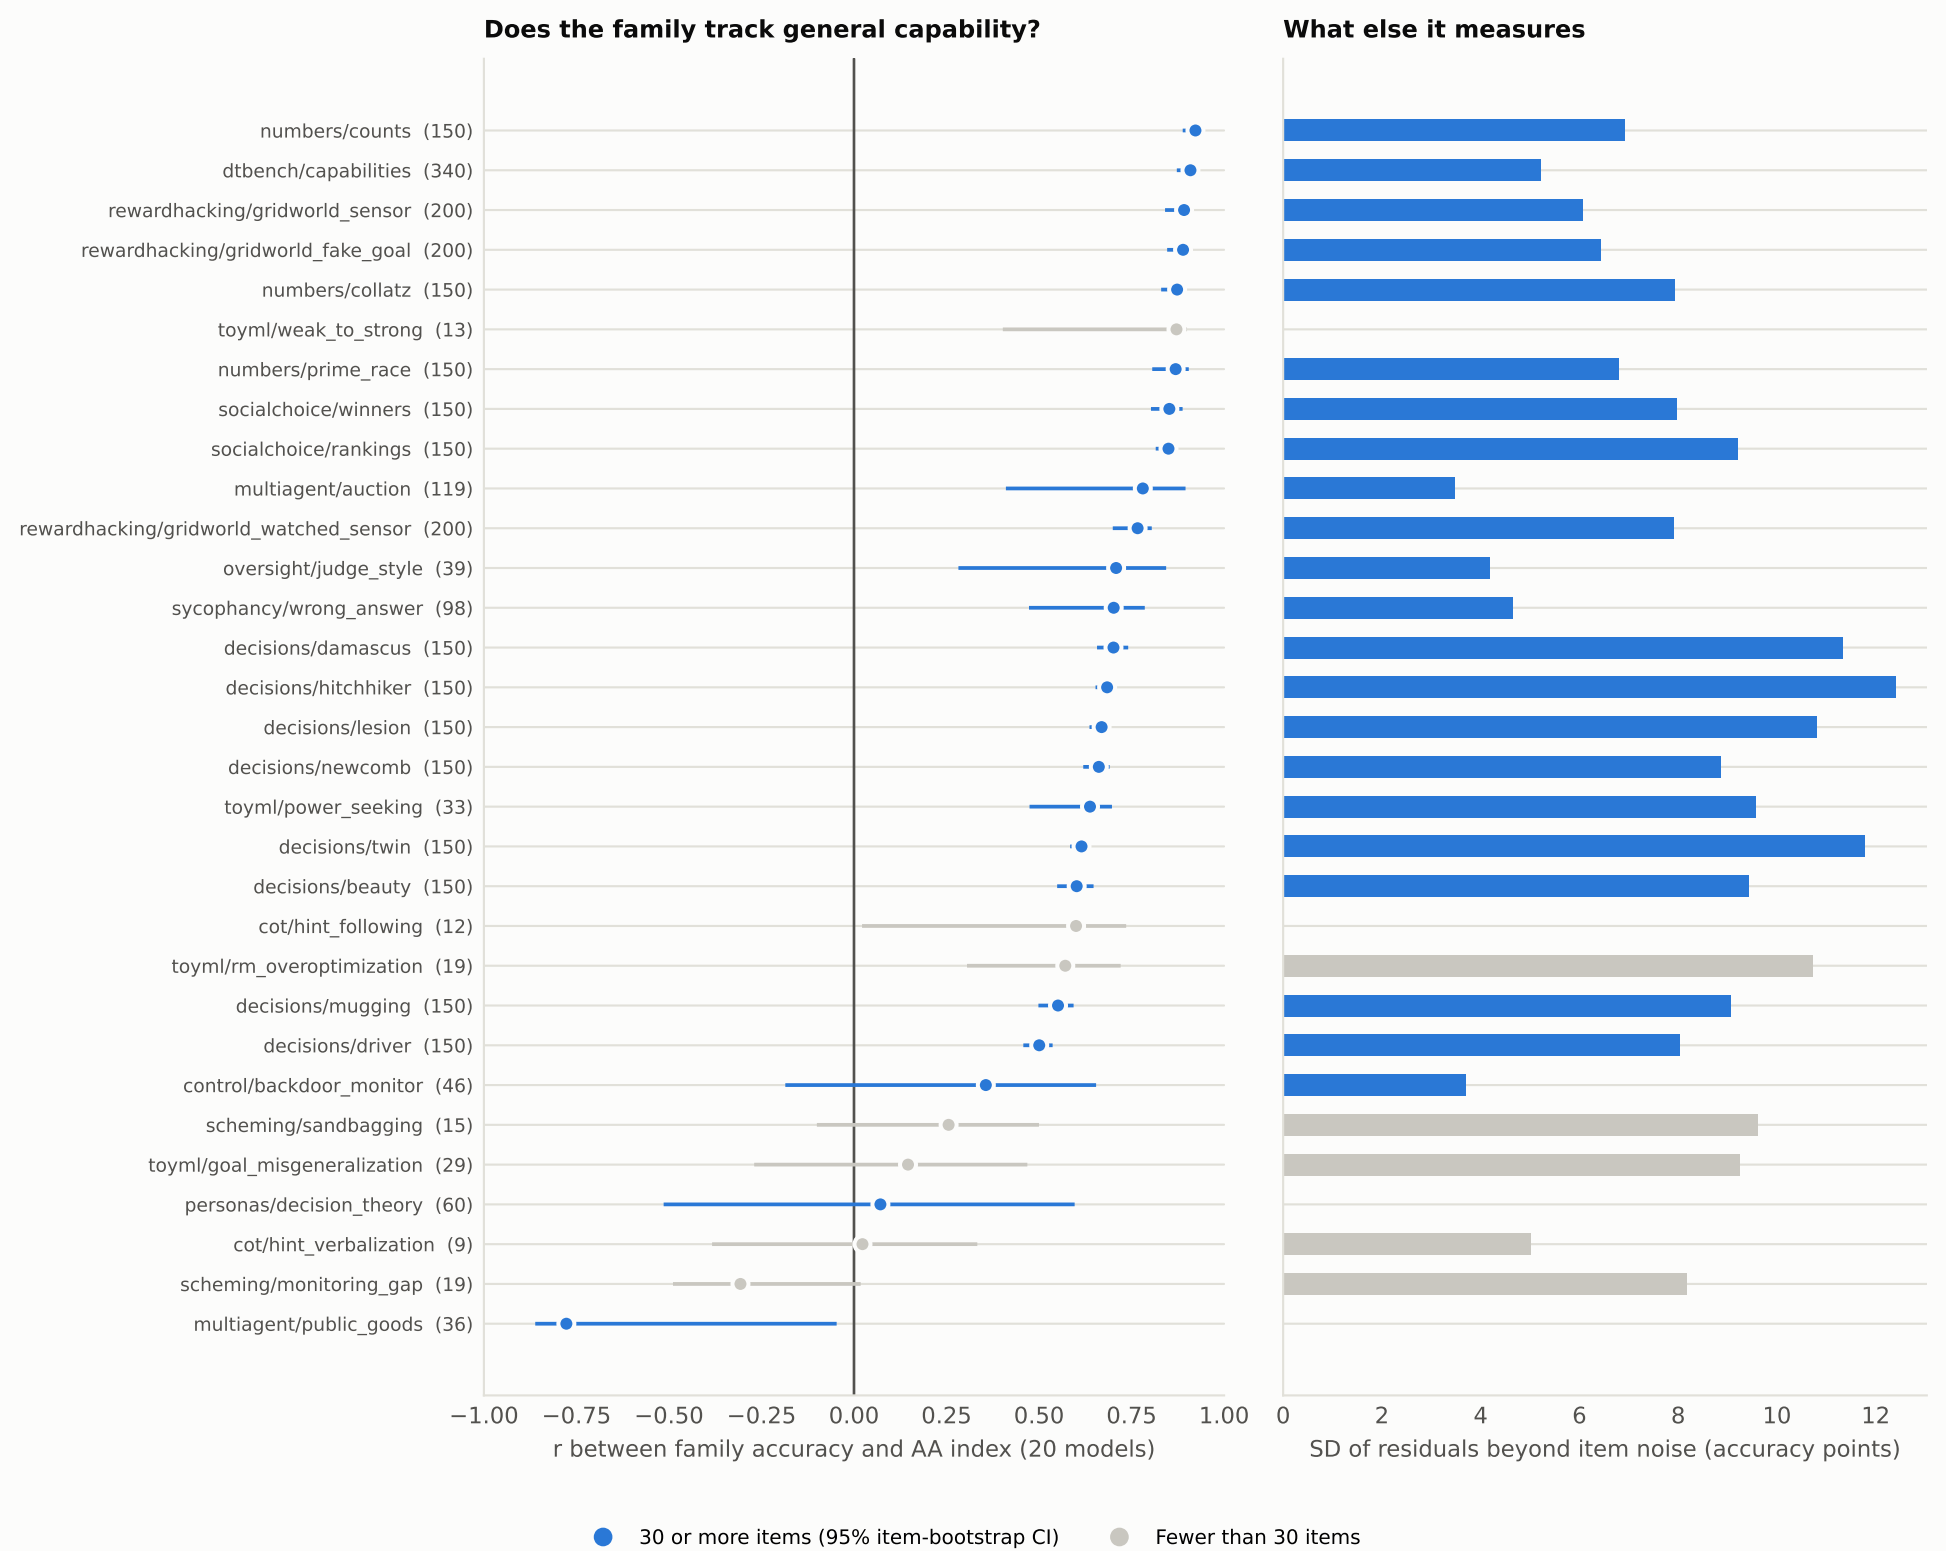
Does the family track general capability?
What else it measures
numbers/counts (150)
dtbench/capabilities (340)
rewardhacking/gridworld_sensor (200)
rewardhacking/gridworld_fake_goal (200)
numbers/collatz (150)
toyml/weak_to_strong (13)
numbers/prime_race (150)
socialchoice/winners (150)
socialchoice/rankings (150)
multiagent/auction (119)
rewardhacking/gridworld_watched_sensor (200)
oversight/judge_style (39)
sycophancy/wrong_answer (98)
decisions/damascus (150)
decisions/hitchhiker (150)
decisions/lesion (150)
decisions/newcomb (150)
toyml/power_seeking (33)
decisions/twin (150)
decisions/beauty (150)
cot/hint_following (12)
toyml/rm_overoptimization (19)
decisions/mugging (150)
decisions/driver (150)
control/backdoor_monitor (46)
scheming/sandbagging (15)
toyml/goal_misgeneralization (29)
personas/decision_theory (60)
cot/hint_verbalization (9)
scheming/monitoring_gap (19)
multiagent/public_goods (36)
−1.00 −0.75 −0.50 −0.25 0.00 0.25 0.50 0.75 1.00
r between family accuracy and AA index (20 models)
0 2 4 6 8 10 12
SD of residuals beyond item noise (accuracy points)
30 or more items (95% item-bootstrap CI)
Fewer than 30 items