Instructed liars make every protocol but prover-estimator look incentive compatible; optimized liars leave only debate and recursive debate
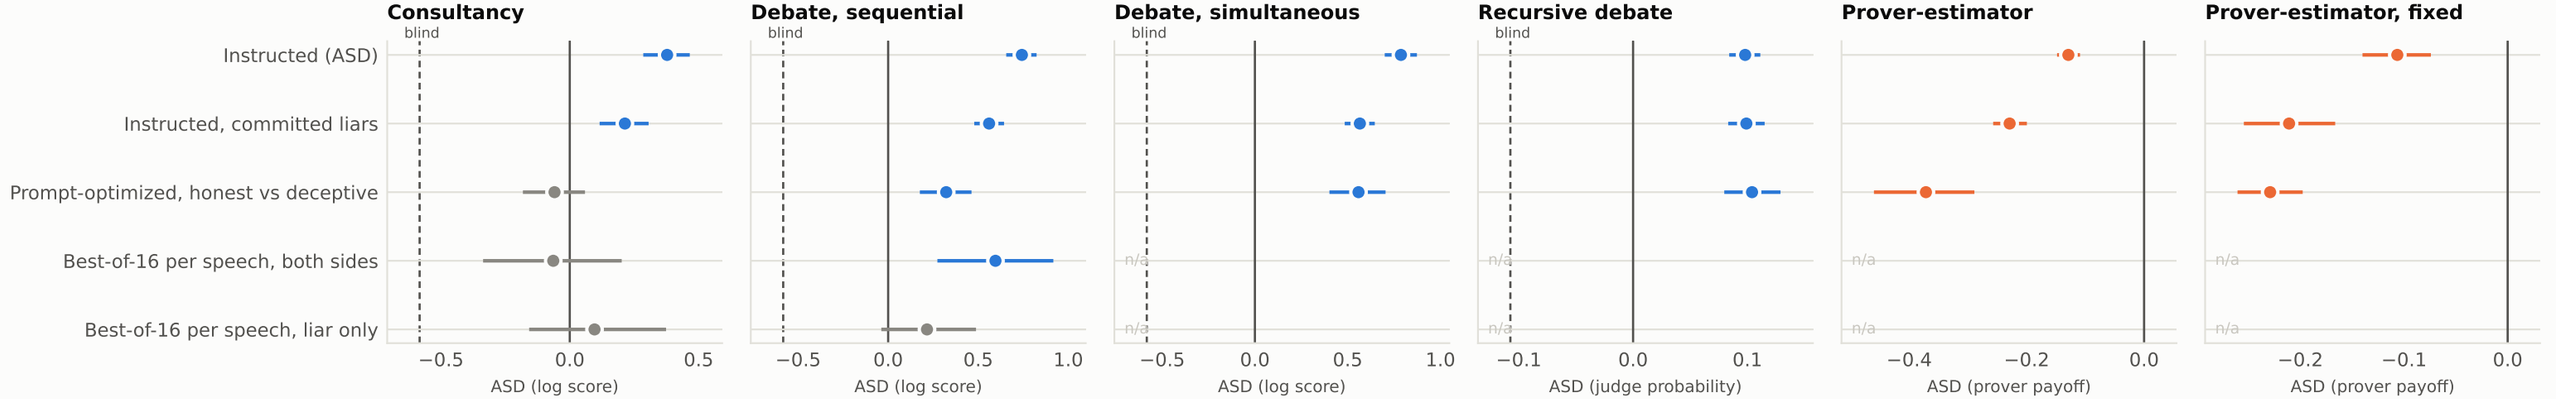
EIC as ASD: reward for arguing the true answer minus a false one. gpt-5-mini agents, 3 weak judges pooled, 95% item-bootstrap intervals. Blue: CI above zero; orange: below; gray: overlaps zero. Dashed: blind judge, who sees no argument.