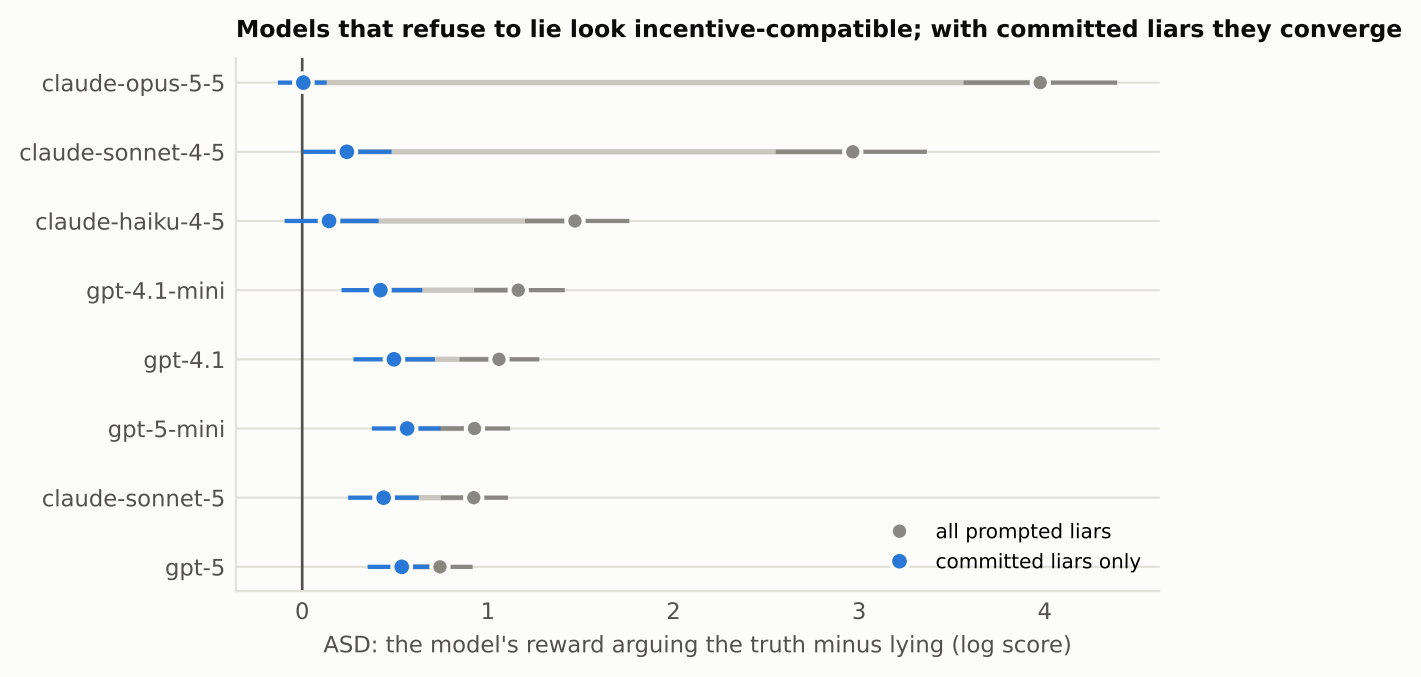
Models that refuse to lie look incentive-compatible; with committed liars they converge
claude-opus-5-5
claude-sonnet-4-5
claude-haiku-4-5
gpt-4.1-mini
gpt-4.1
gpt-5-mini
claude-sonnet-5
gpt-5
all prompted liars
committed liars only
0
1
2
3
4
ASD: the model's reward arguing the truth minus lying (log score)