Self-play: honesty pays at equilibrium in simultaneous debate; sequential debate's incentive is optimized away
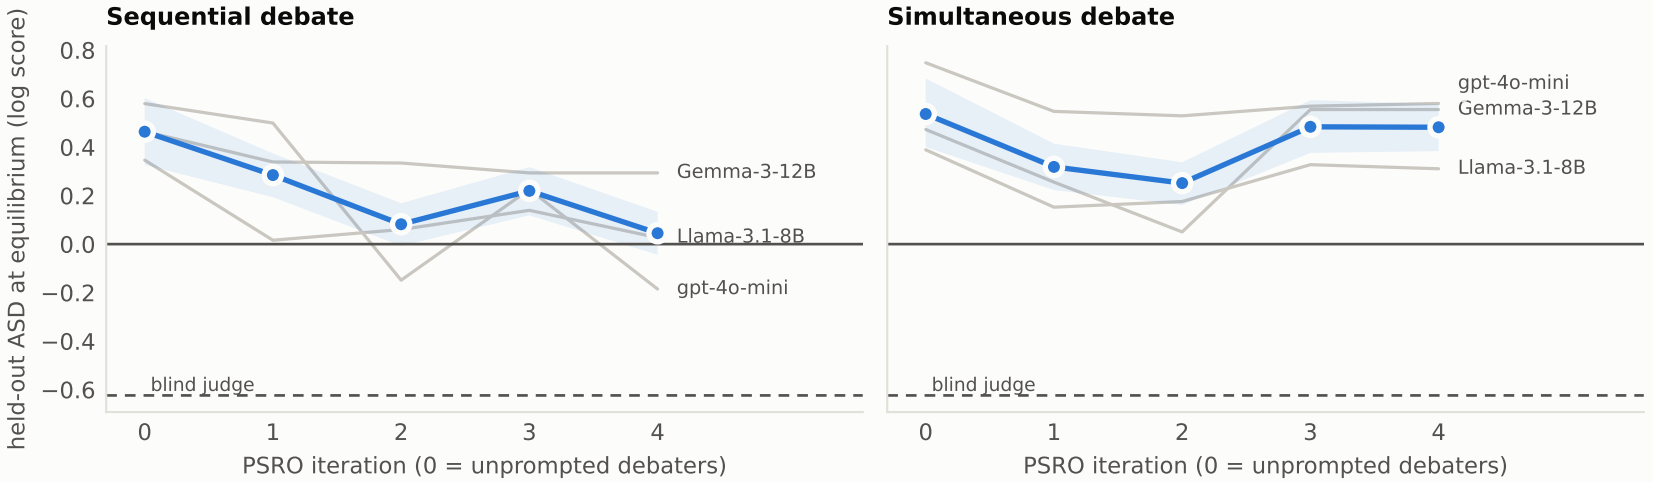
Honest and lying gpt-5-mini strategies best-respond to each other's Nash mixture (prompt search on 100 training items, judge = mean of 3 samples); ASD = E[honest - liar reward] at each equilibrium on 390 held-out items.
Blue: judges pooled, 95% item-bootstrap band; gray: each judge.