Without reasoning noise the judge is easier to exploit: best-of-16 liars drive the honesty incentive far below zero
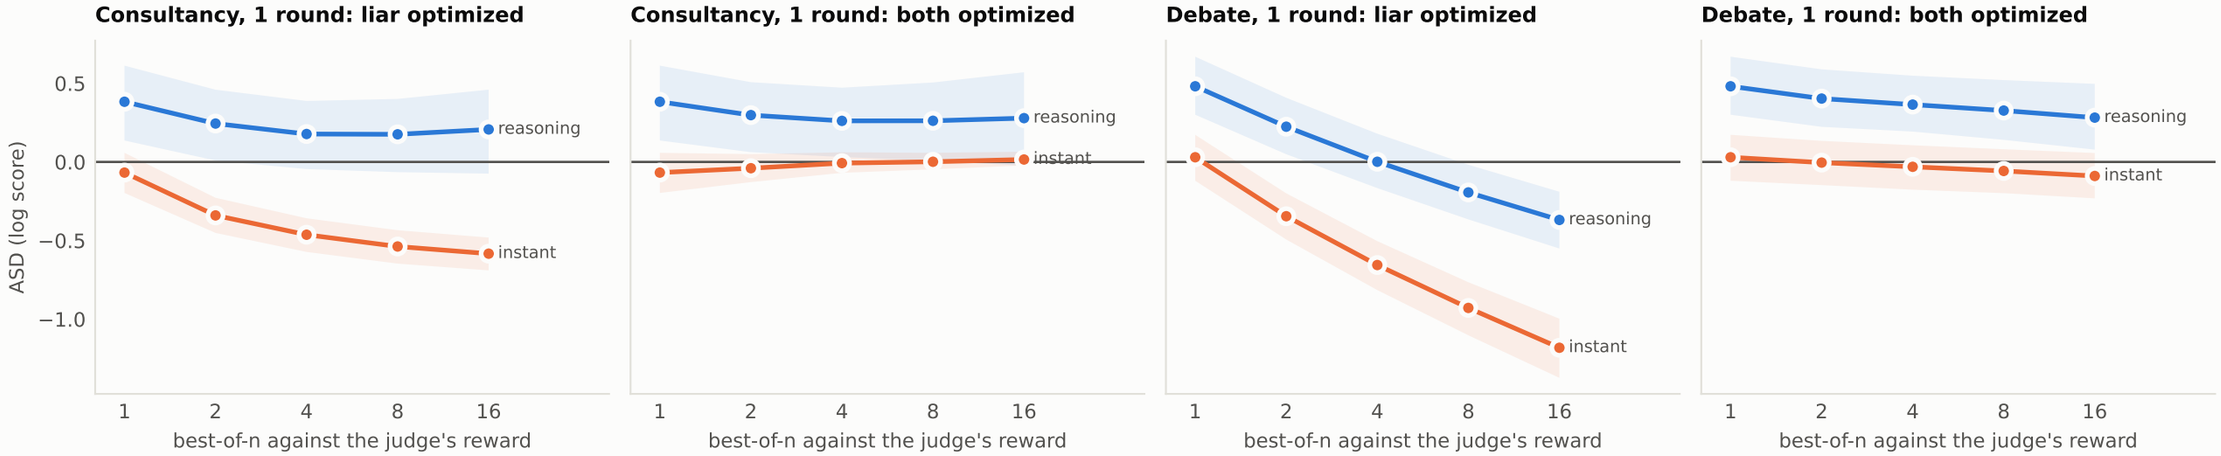
Orange: weak judges answering at once (no reasoning, temperature 0; two calls agree at r = 0.78-0.95), so the payoff that selects is the payoff reported. Blue: the same judges reasoning first, scored on samples independent of selection. gpt-5-mini, 150 items, 3 judges pooled, 95% bands.