Debate edges above the judges with no argument, more so for smarter agents; consultancy does not
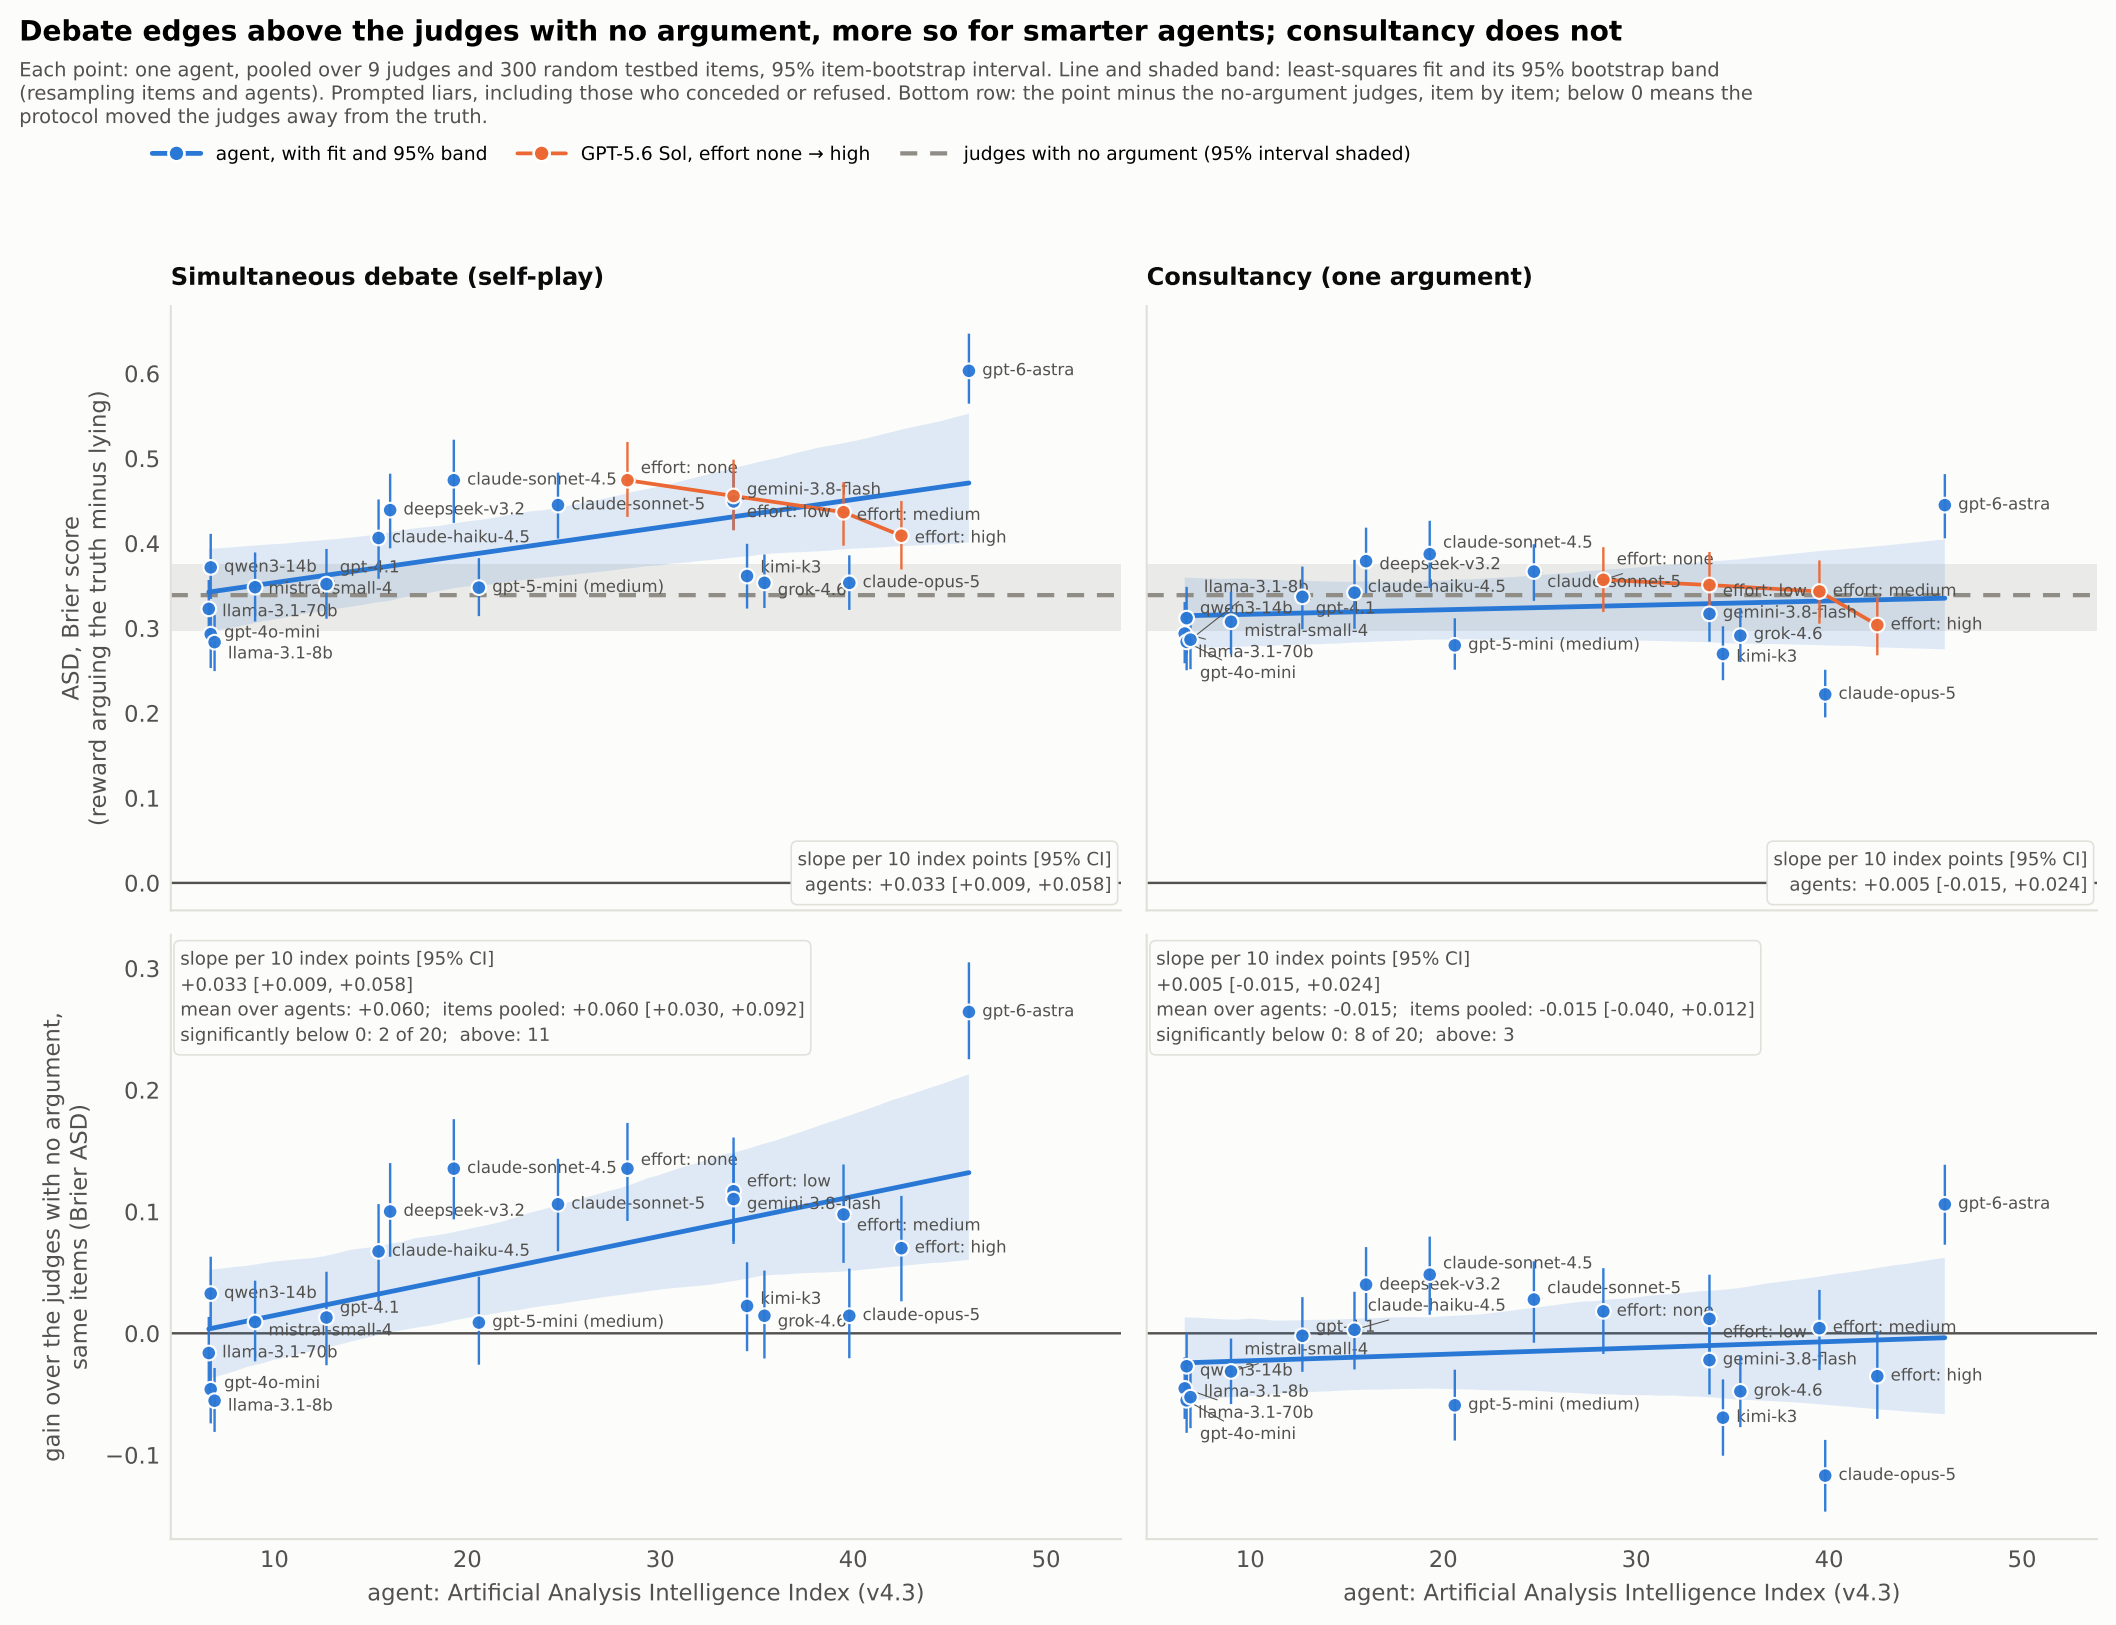
Each point: one agent, pooled over 9 judges and 300 random testbed items, 95% item-bootstrap interval. Line and shaded band: least-squares fit and its 95% bootstrap band (resampling items and agents). Prompted liars, including those who conceded or refused. Bottom row: the point minus the no-argument judges, item by item; below 0 means the protocol moved the judges away from the truth.