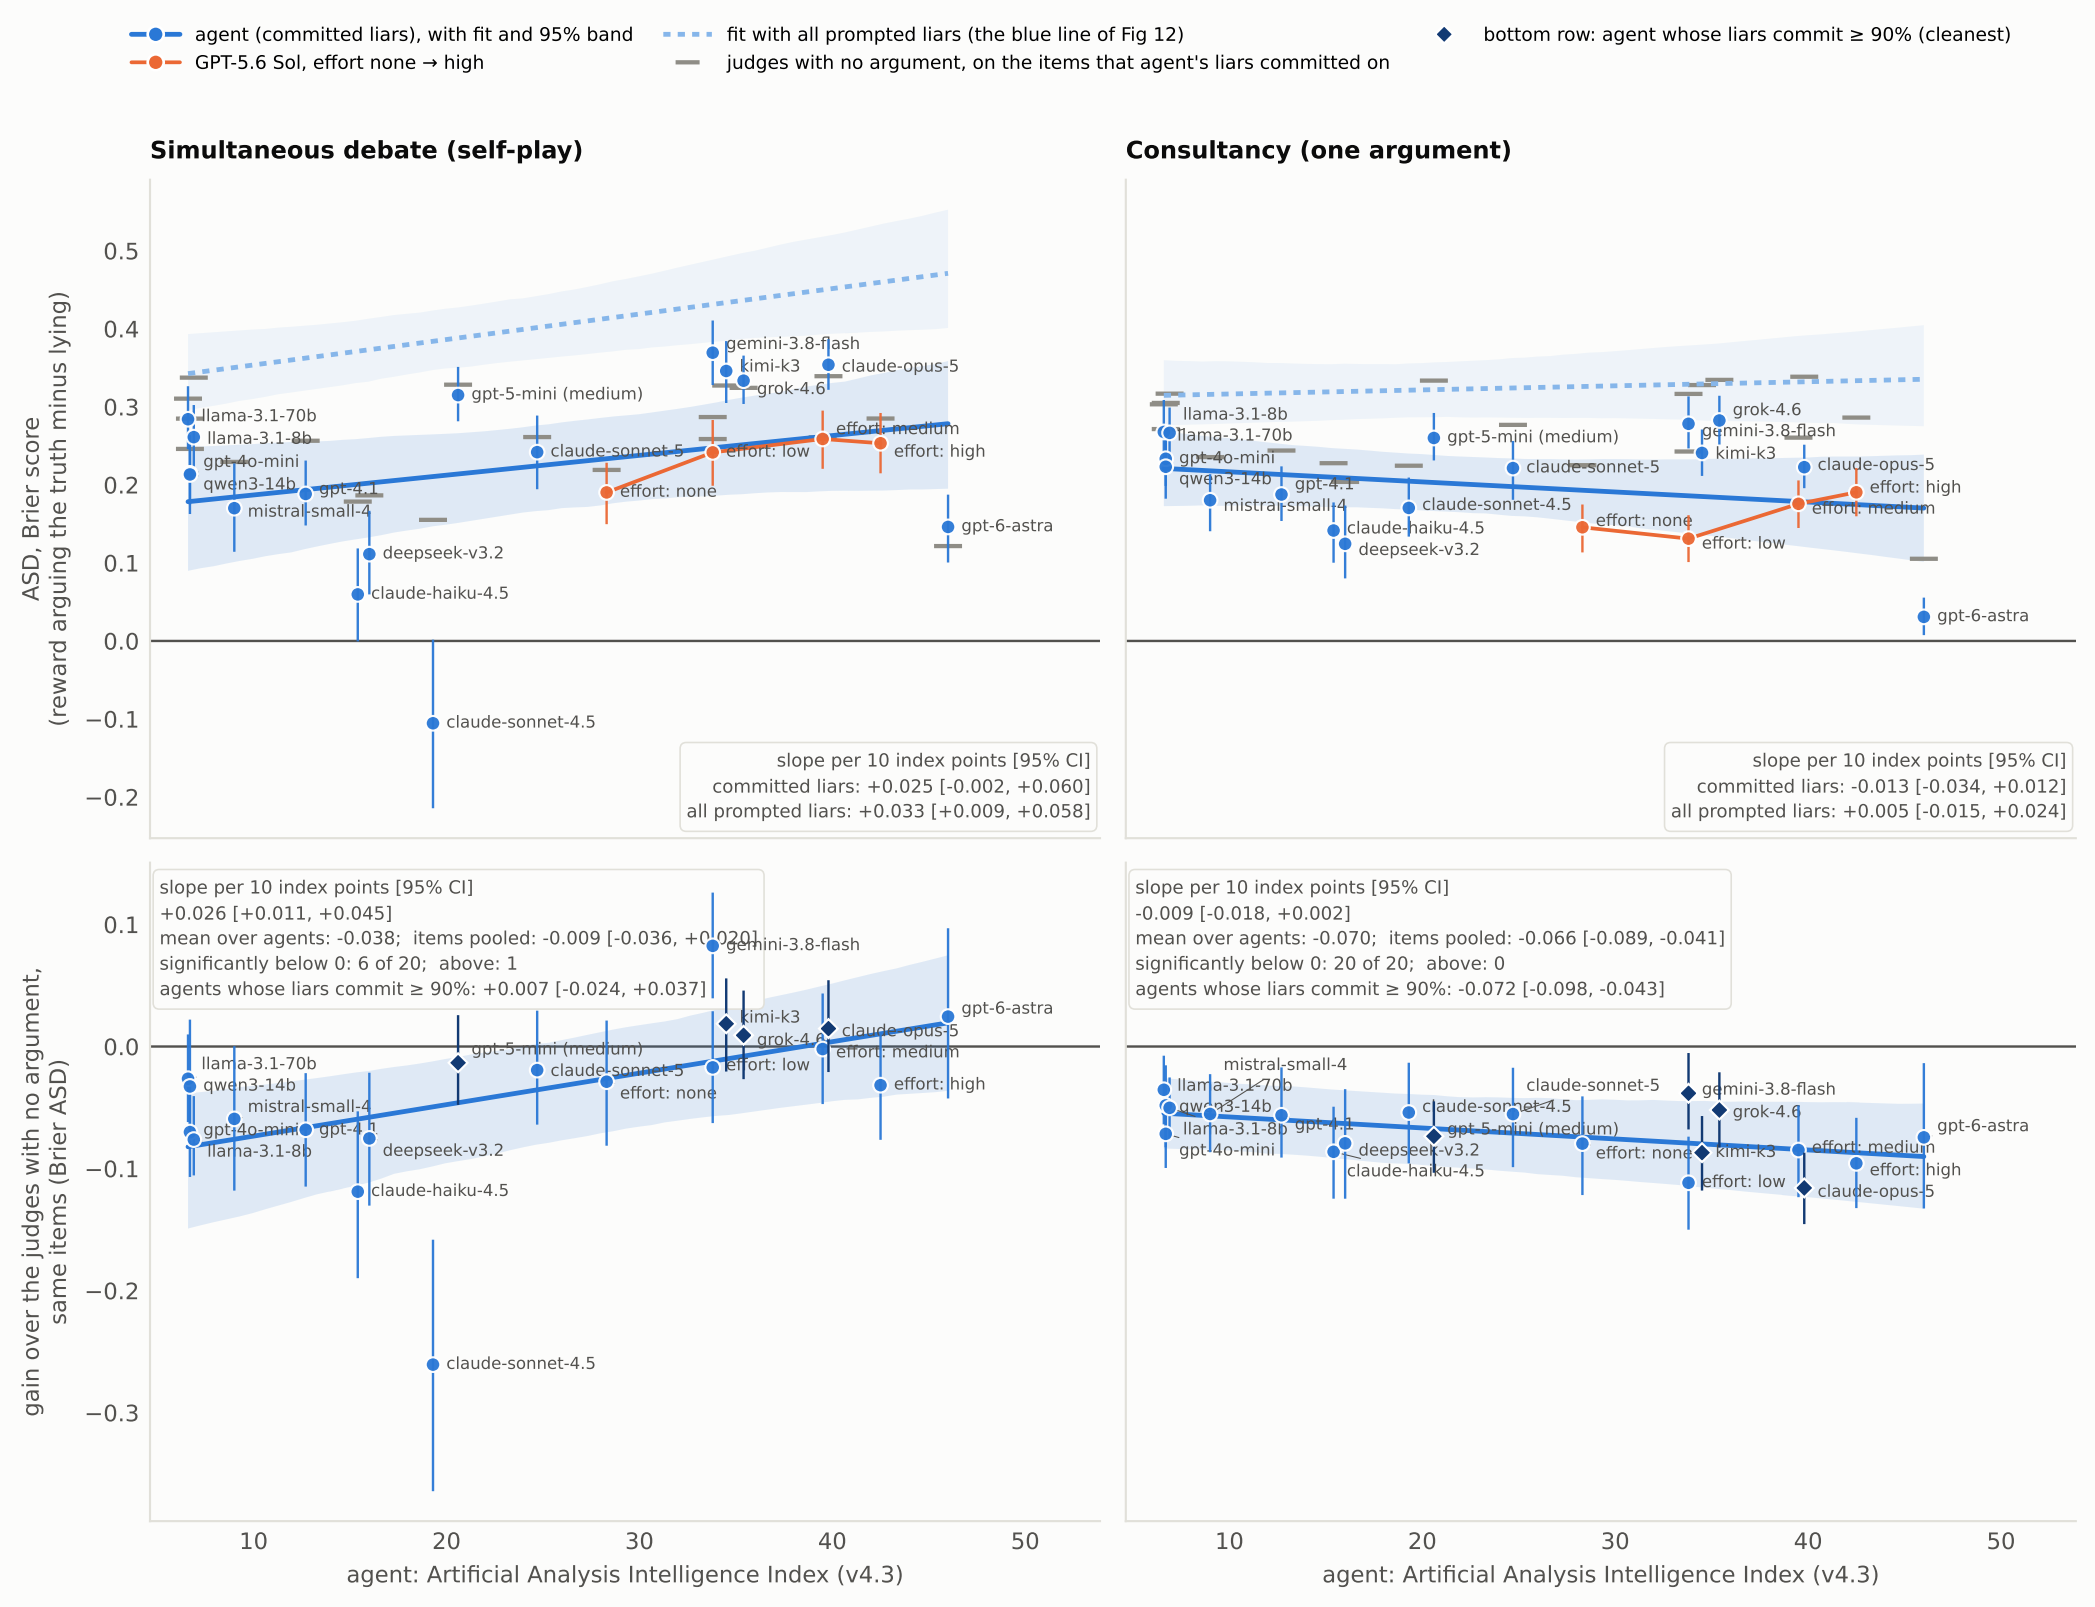
agent (committed liars), with fit and 95% band
GPT-5.6 Sol, effort none → high
fit with all prompted liars (the blue line of Fig 12)
judges with no argument, on the items that agent's liars committed on
bottom row: agent whose liars commit ≥ 90% (cleanest)
Simultaneous debate (self-play)
Consultancy (one argument)
ASD, Brier score
(reward arguing the truth minus lying)
gain over the judges with no argument,
same items (Brier ASD)
agent: Artificial Analysis Intelligence Index (v4.3)
slope per 10 index points [95% CI]
committed liars: +0.025 [-0.002, +0.060]
all prompted liars: +0.033 [+0.009, +0.058]
slope per 10 index points [95% CI]
committed liars: -0.013 [-0.034, +0.012]
all prompted liars: +0.005 [-0.015, +0.024]
slope per 10 index points [95% CI]
+0.026 [+0.011, +0.045]
mean over agents: -0.038; items pooled: -0.009 [-0.036, +0.020]
significantly below 0: 6 of 20; above: 1
agents whose liars commit ≥ 90%: +0.007 [-0.024, +0.037]
slope per 10 index points [95% CI]
-0.009 [-0.018, +0.002]
mean over agents: -0.070; items pooled: -0.066 [-0.089, -0.041]
significantly below 0: 20 of 20; above: 0
agents whose liars commit ≥ 90%: -0.072 [-0.098, -0.043]
llama-3.1-8b
llama-3.1-70b
gpt-4o-mini
qwen3-14b
gpt-4.1
mistral-small-4
deepseek-v3.2
claude-haiku-4.5
claude-sonnet-4.5
gpt-5-mini (medium)
gemini-3.8-flash
kimi-k3
grok-4.6
claude-opus-5
gpt-6-astra
effort: none
effort: low
effort: medium
effort: high