Stronger judges reward honesty more, but no more than they do with no argument at all

Each point: one judge, pooled over 20 agent configurations and 300 random testbed items, 95% item-bootstrap interval. Bottom row: the protocol's ASD minus the same judge's ASD with no argument, item by item.
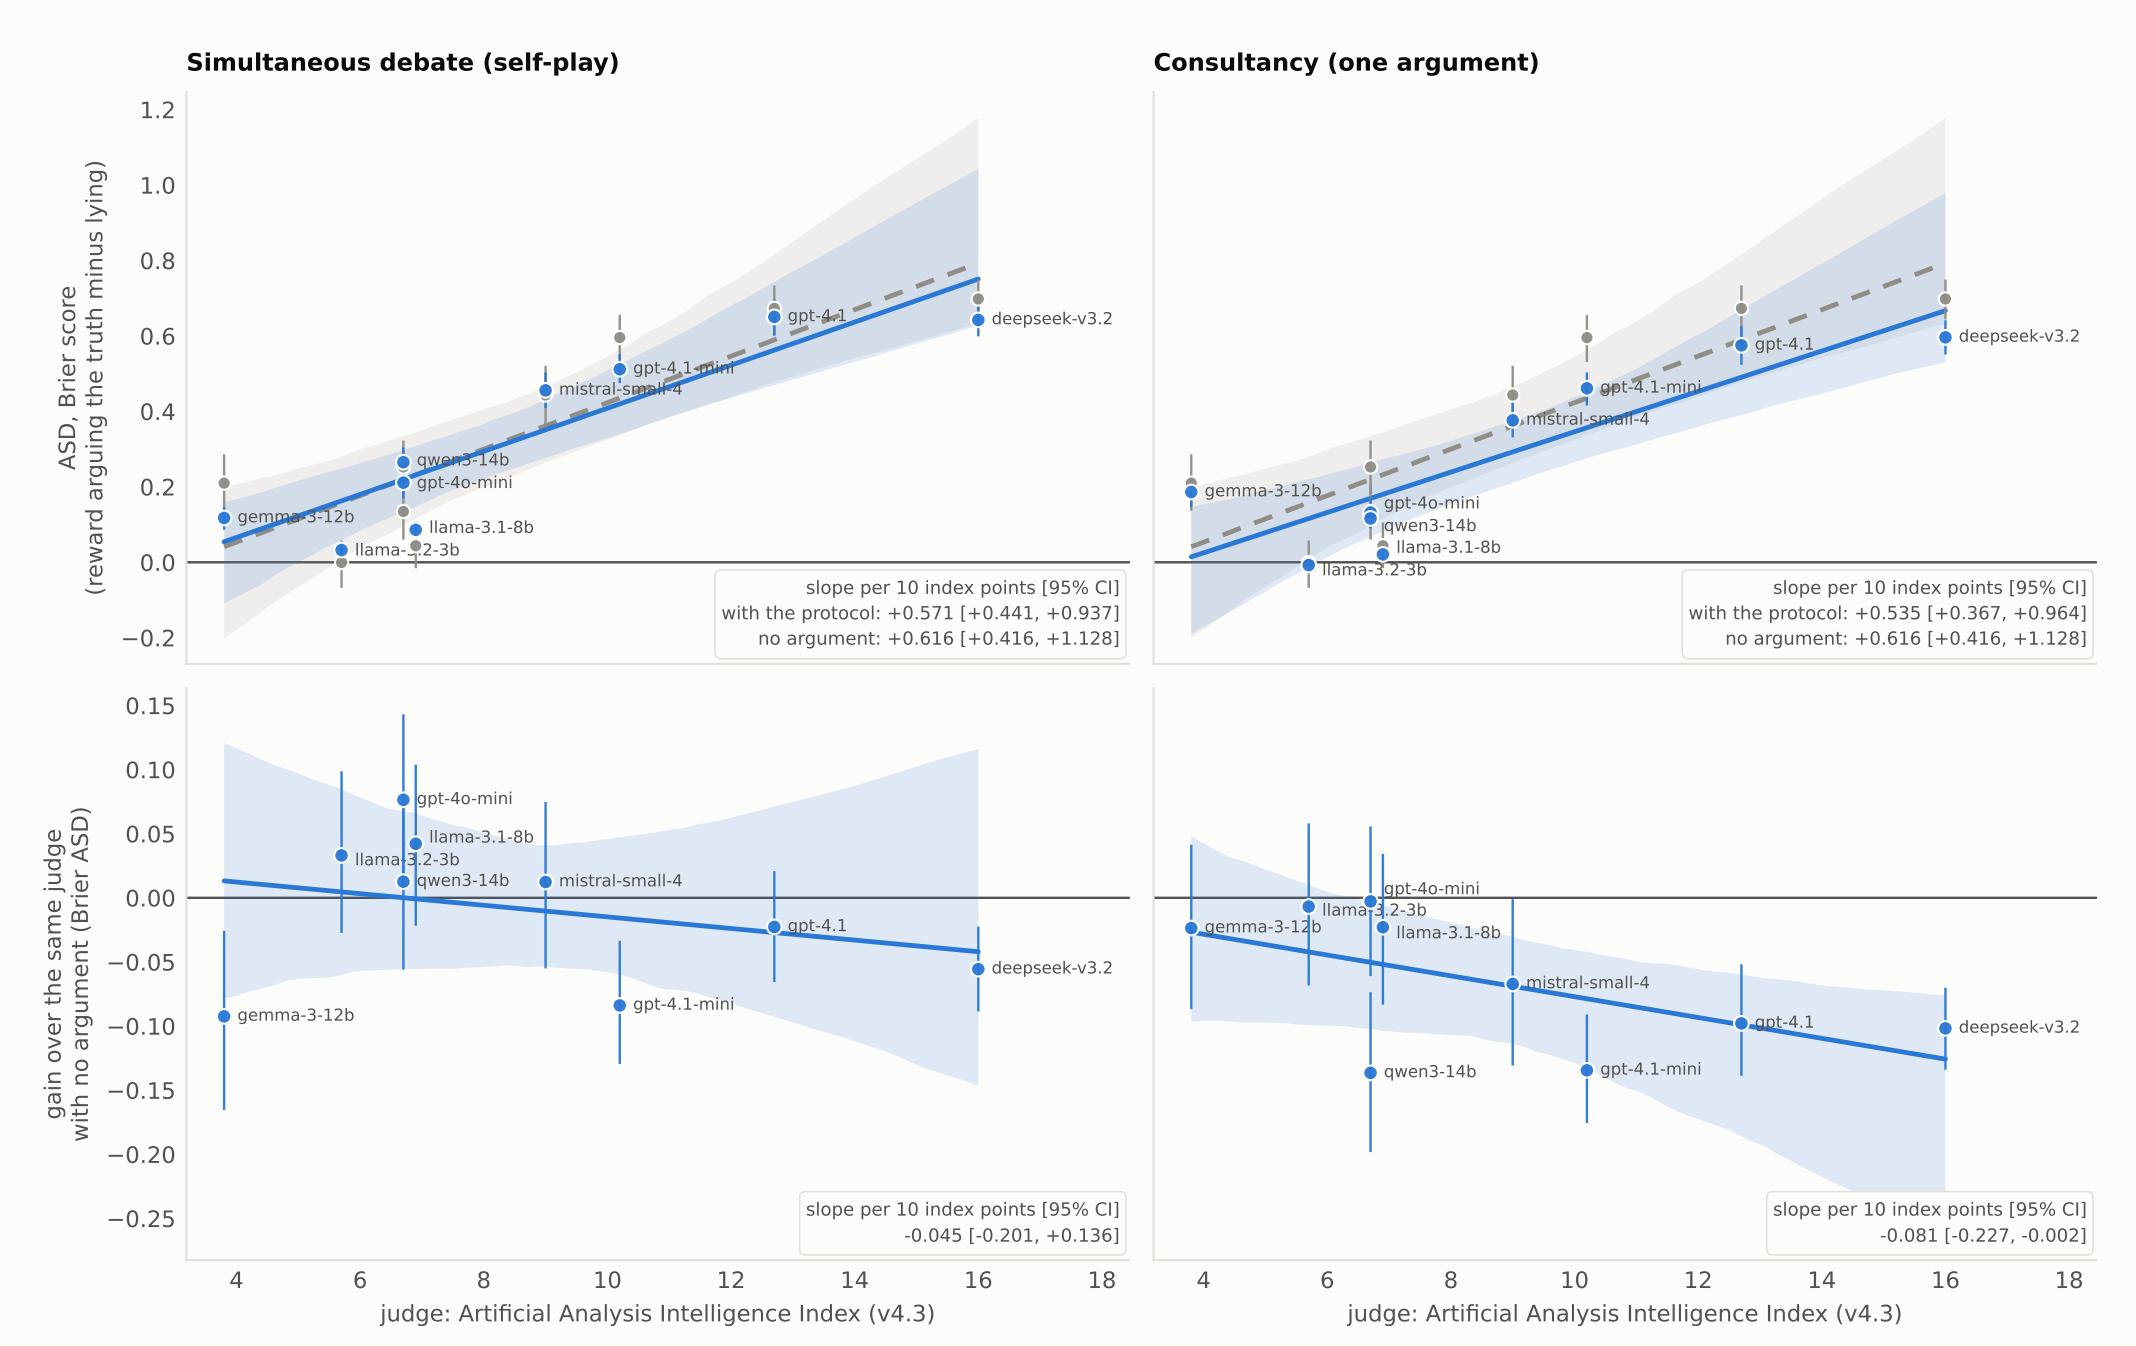
judge with the protocol (committed liars) | same judge, no argument (blind baseline) | 95% bootstrap band of the fit (items and judges)