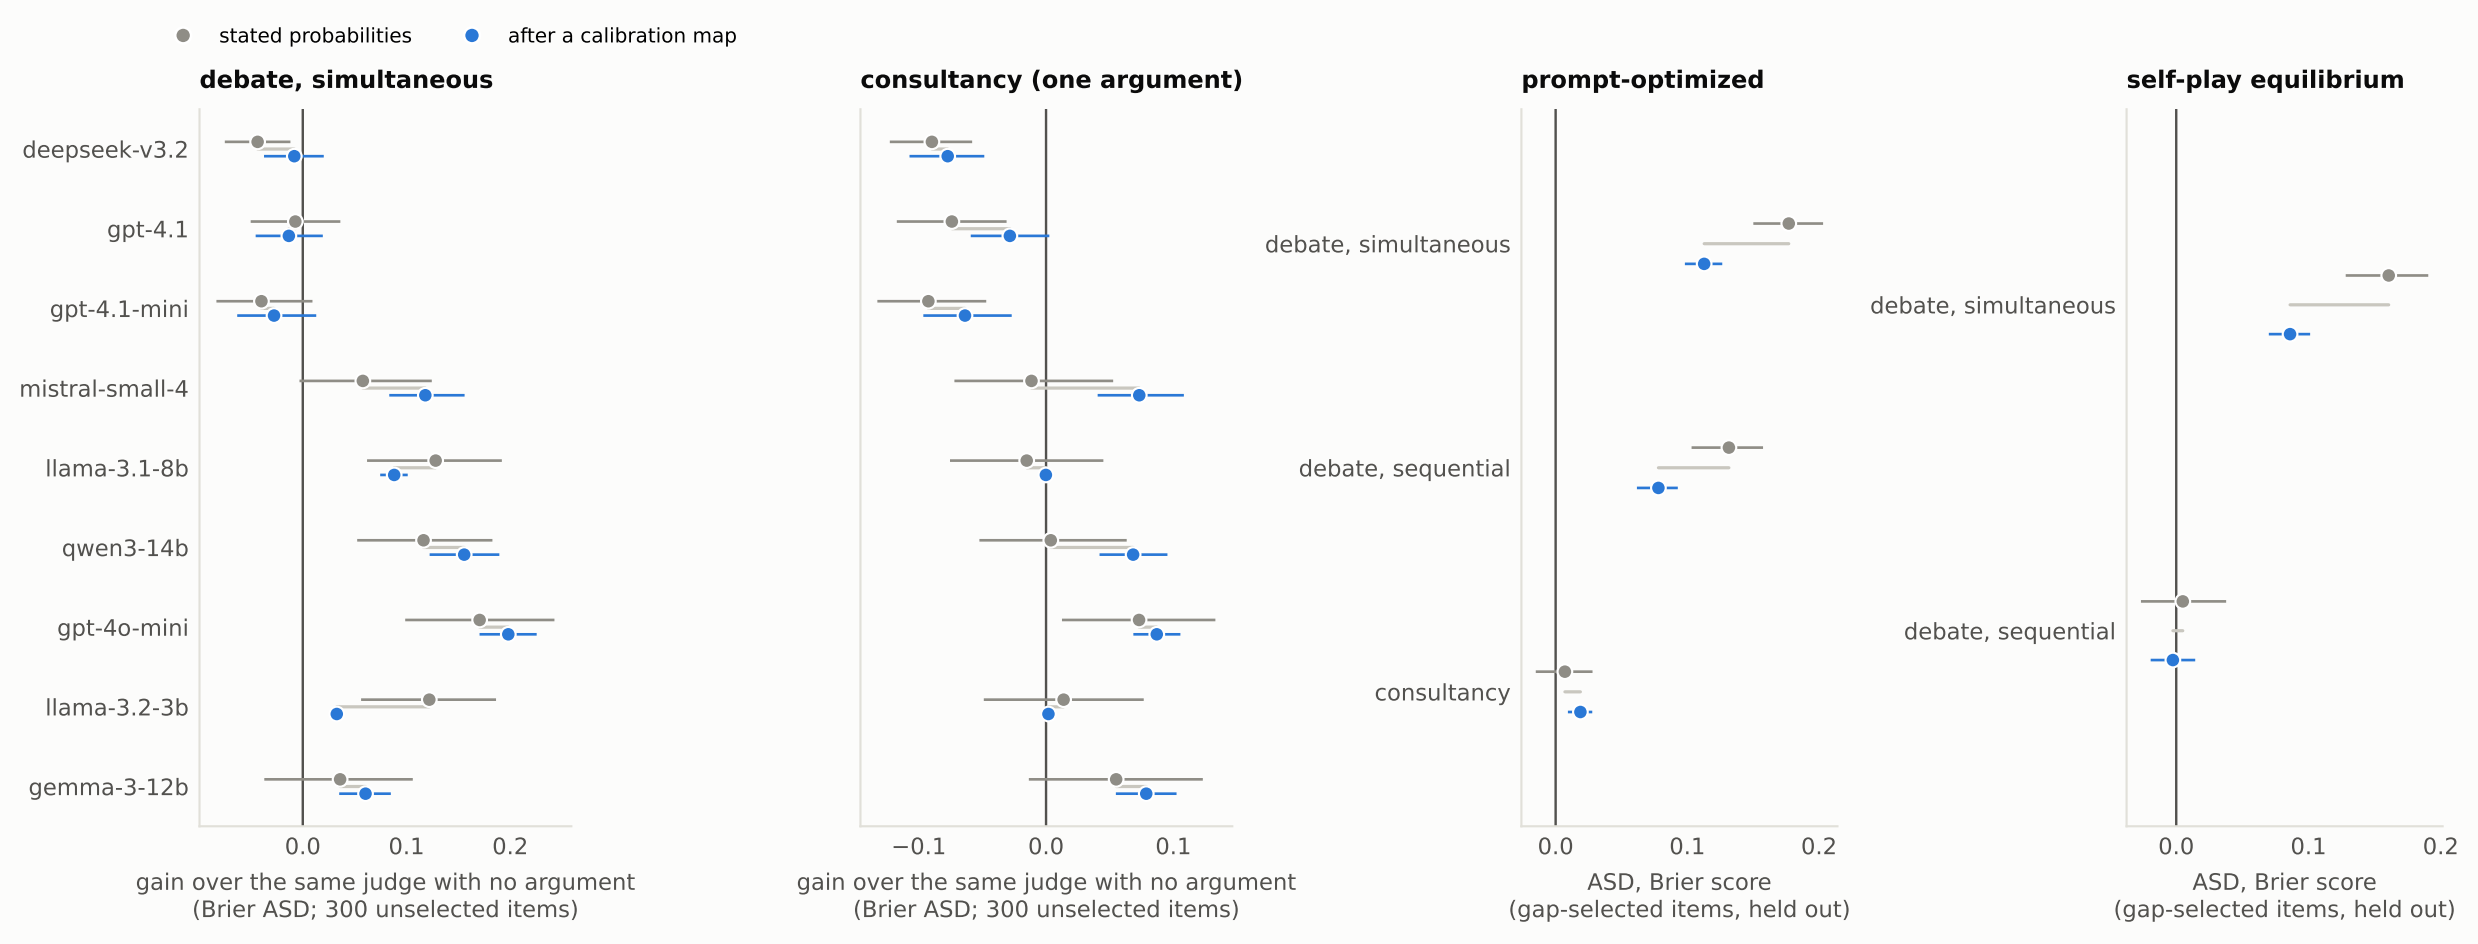
stated probabilities
after a calibration map
debate, simultaneous
consultancy (one argument)
prompt-optimized
self-play equilibrium
deepseek-v3.2
gpt-4.1
gpt-4.1-mini
mistral-small-4
llama-3.1-8b
qwen3-14b
gpt-4o-mini
llama-3.2-3b
gemma-3-12b
debate, simultaneous
debate, sequential
consultancy
0.0
0.1
0.2
−0.1
gain over the same judge with no argument
(Brier ASD; 300 unselected items)
ASD, Brier score
(gap-selected items, held out)