An optimized liar erases most of debate's honesty incentive, and smarter agents don't lose more of it
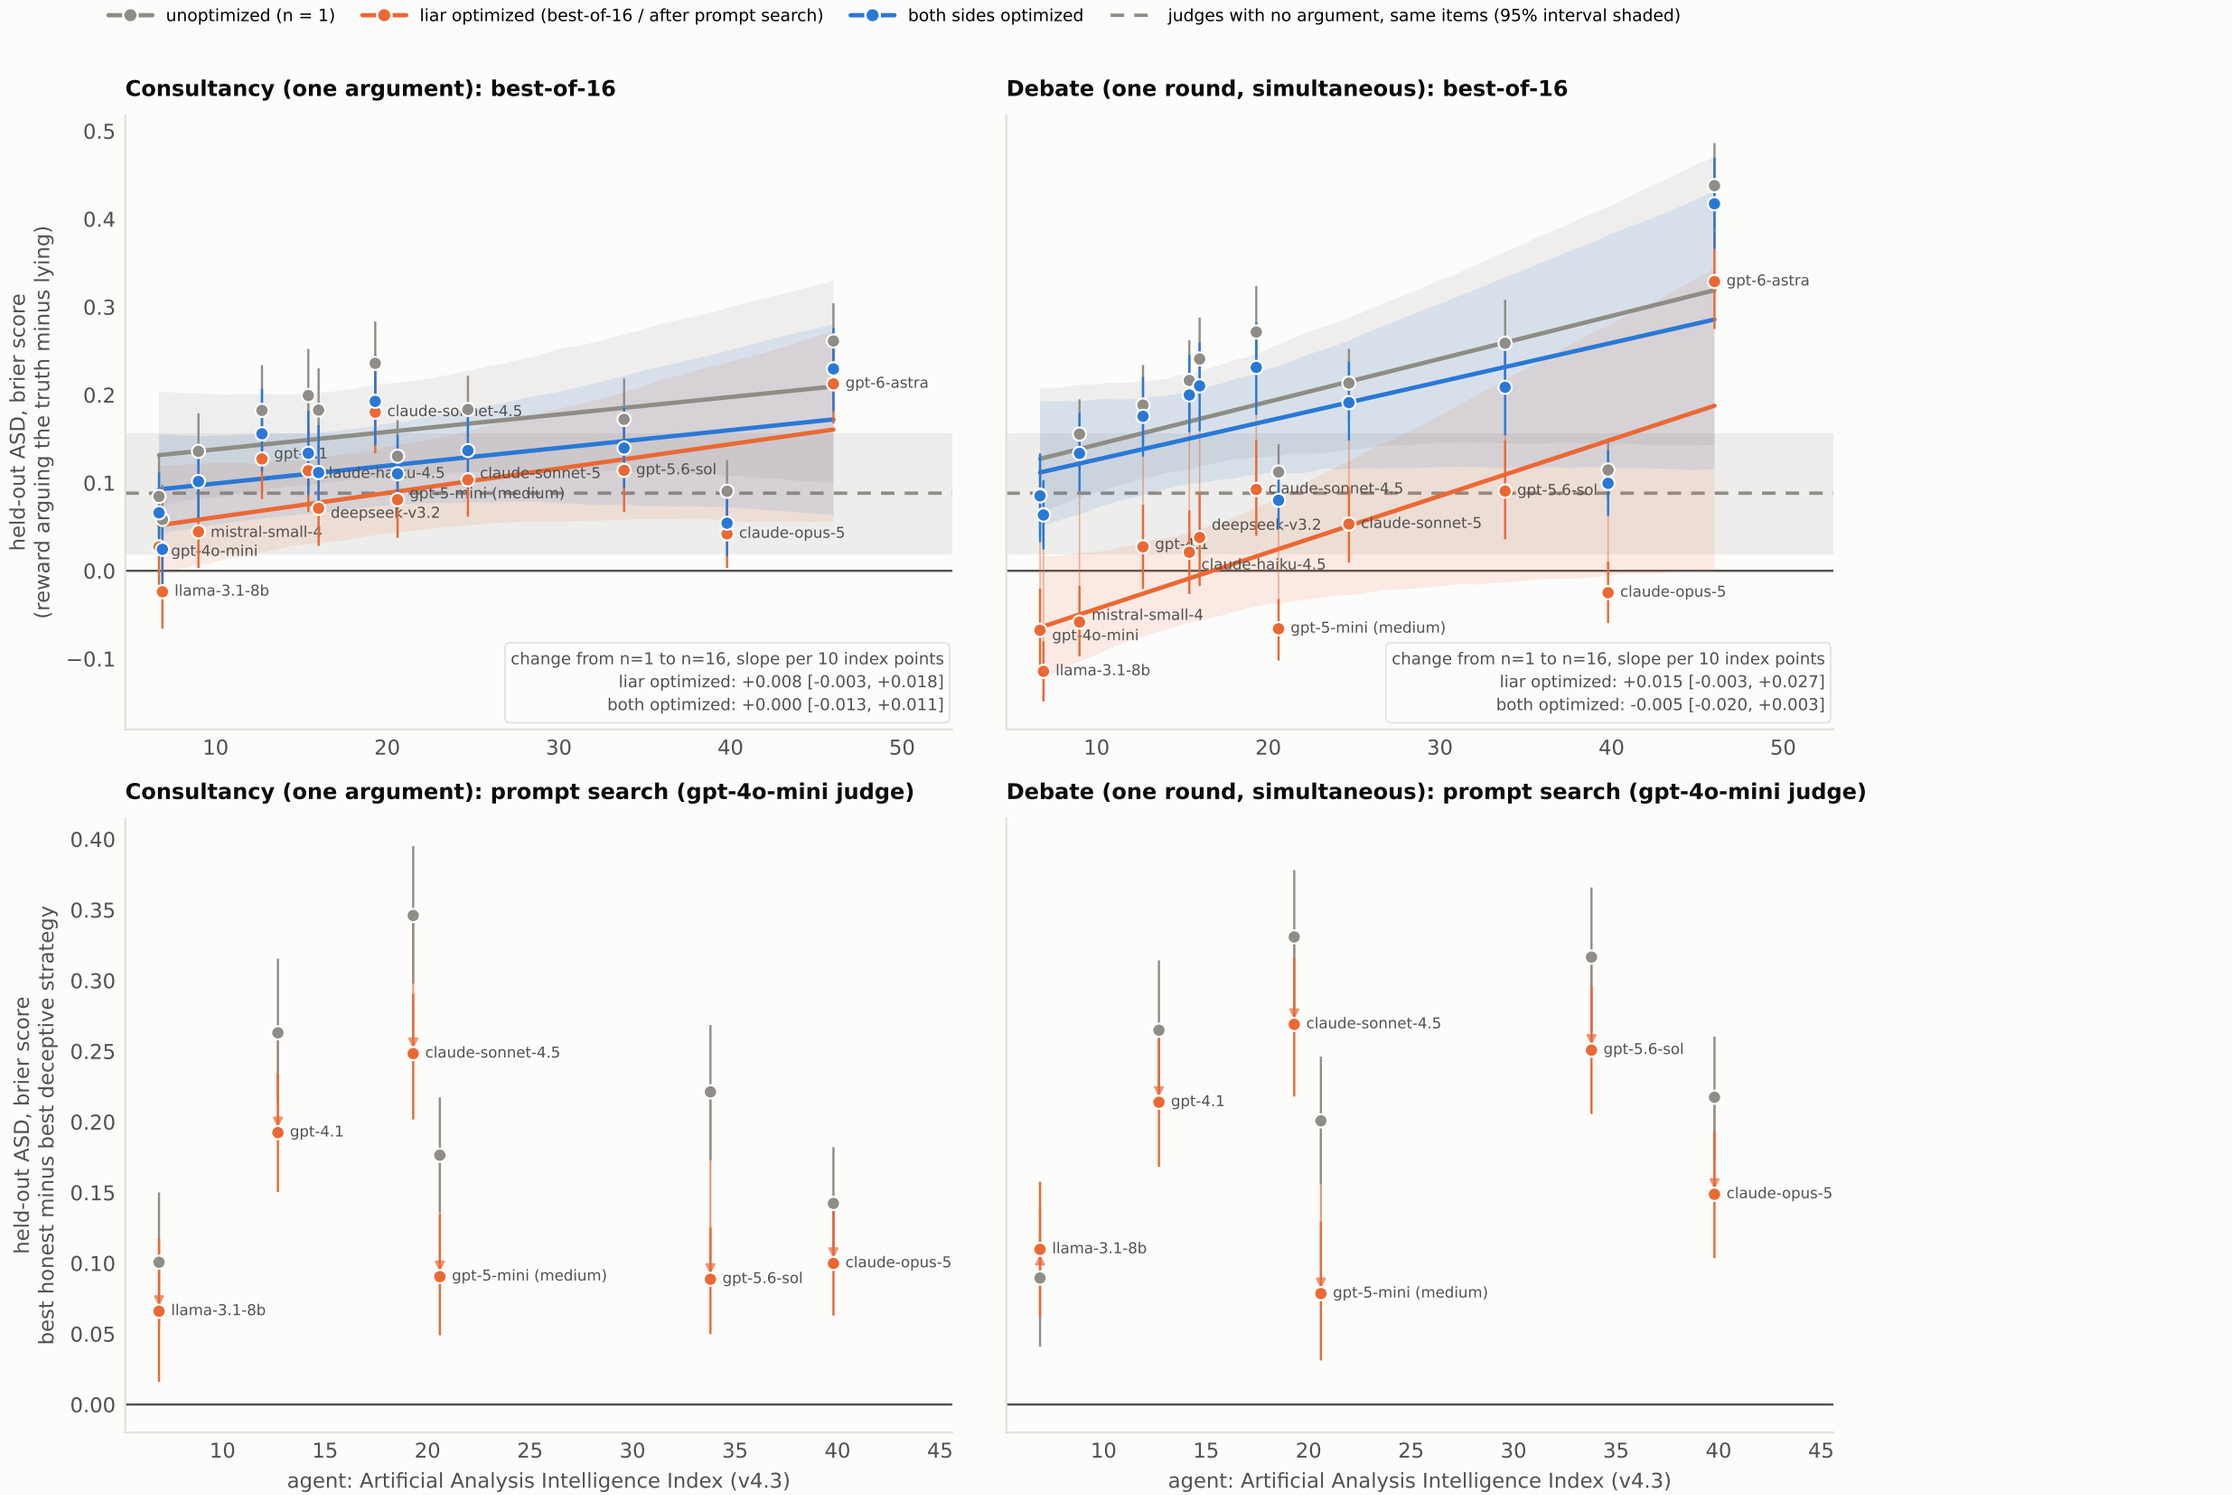
Top: each point is one agent, pooled over 3 weak judges, scored on independent judge samples; lines and bands: least-squares fits with 95% bootstrap bands (items and agents). Bottom: arrows from the unprompted agent to the end of the prompt search.