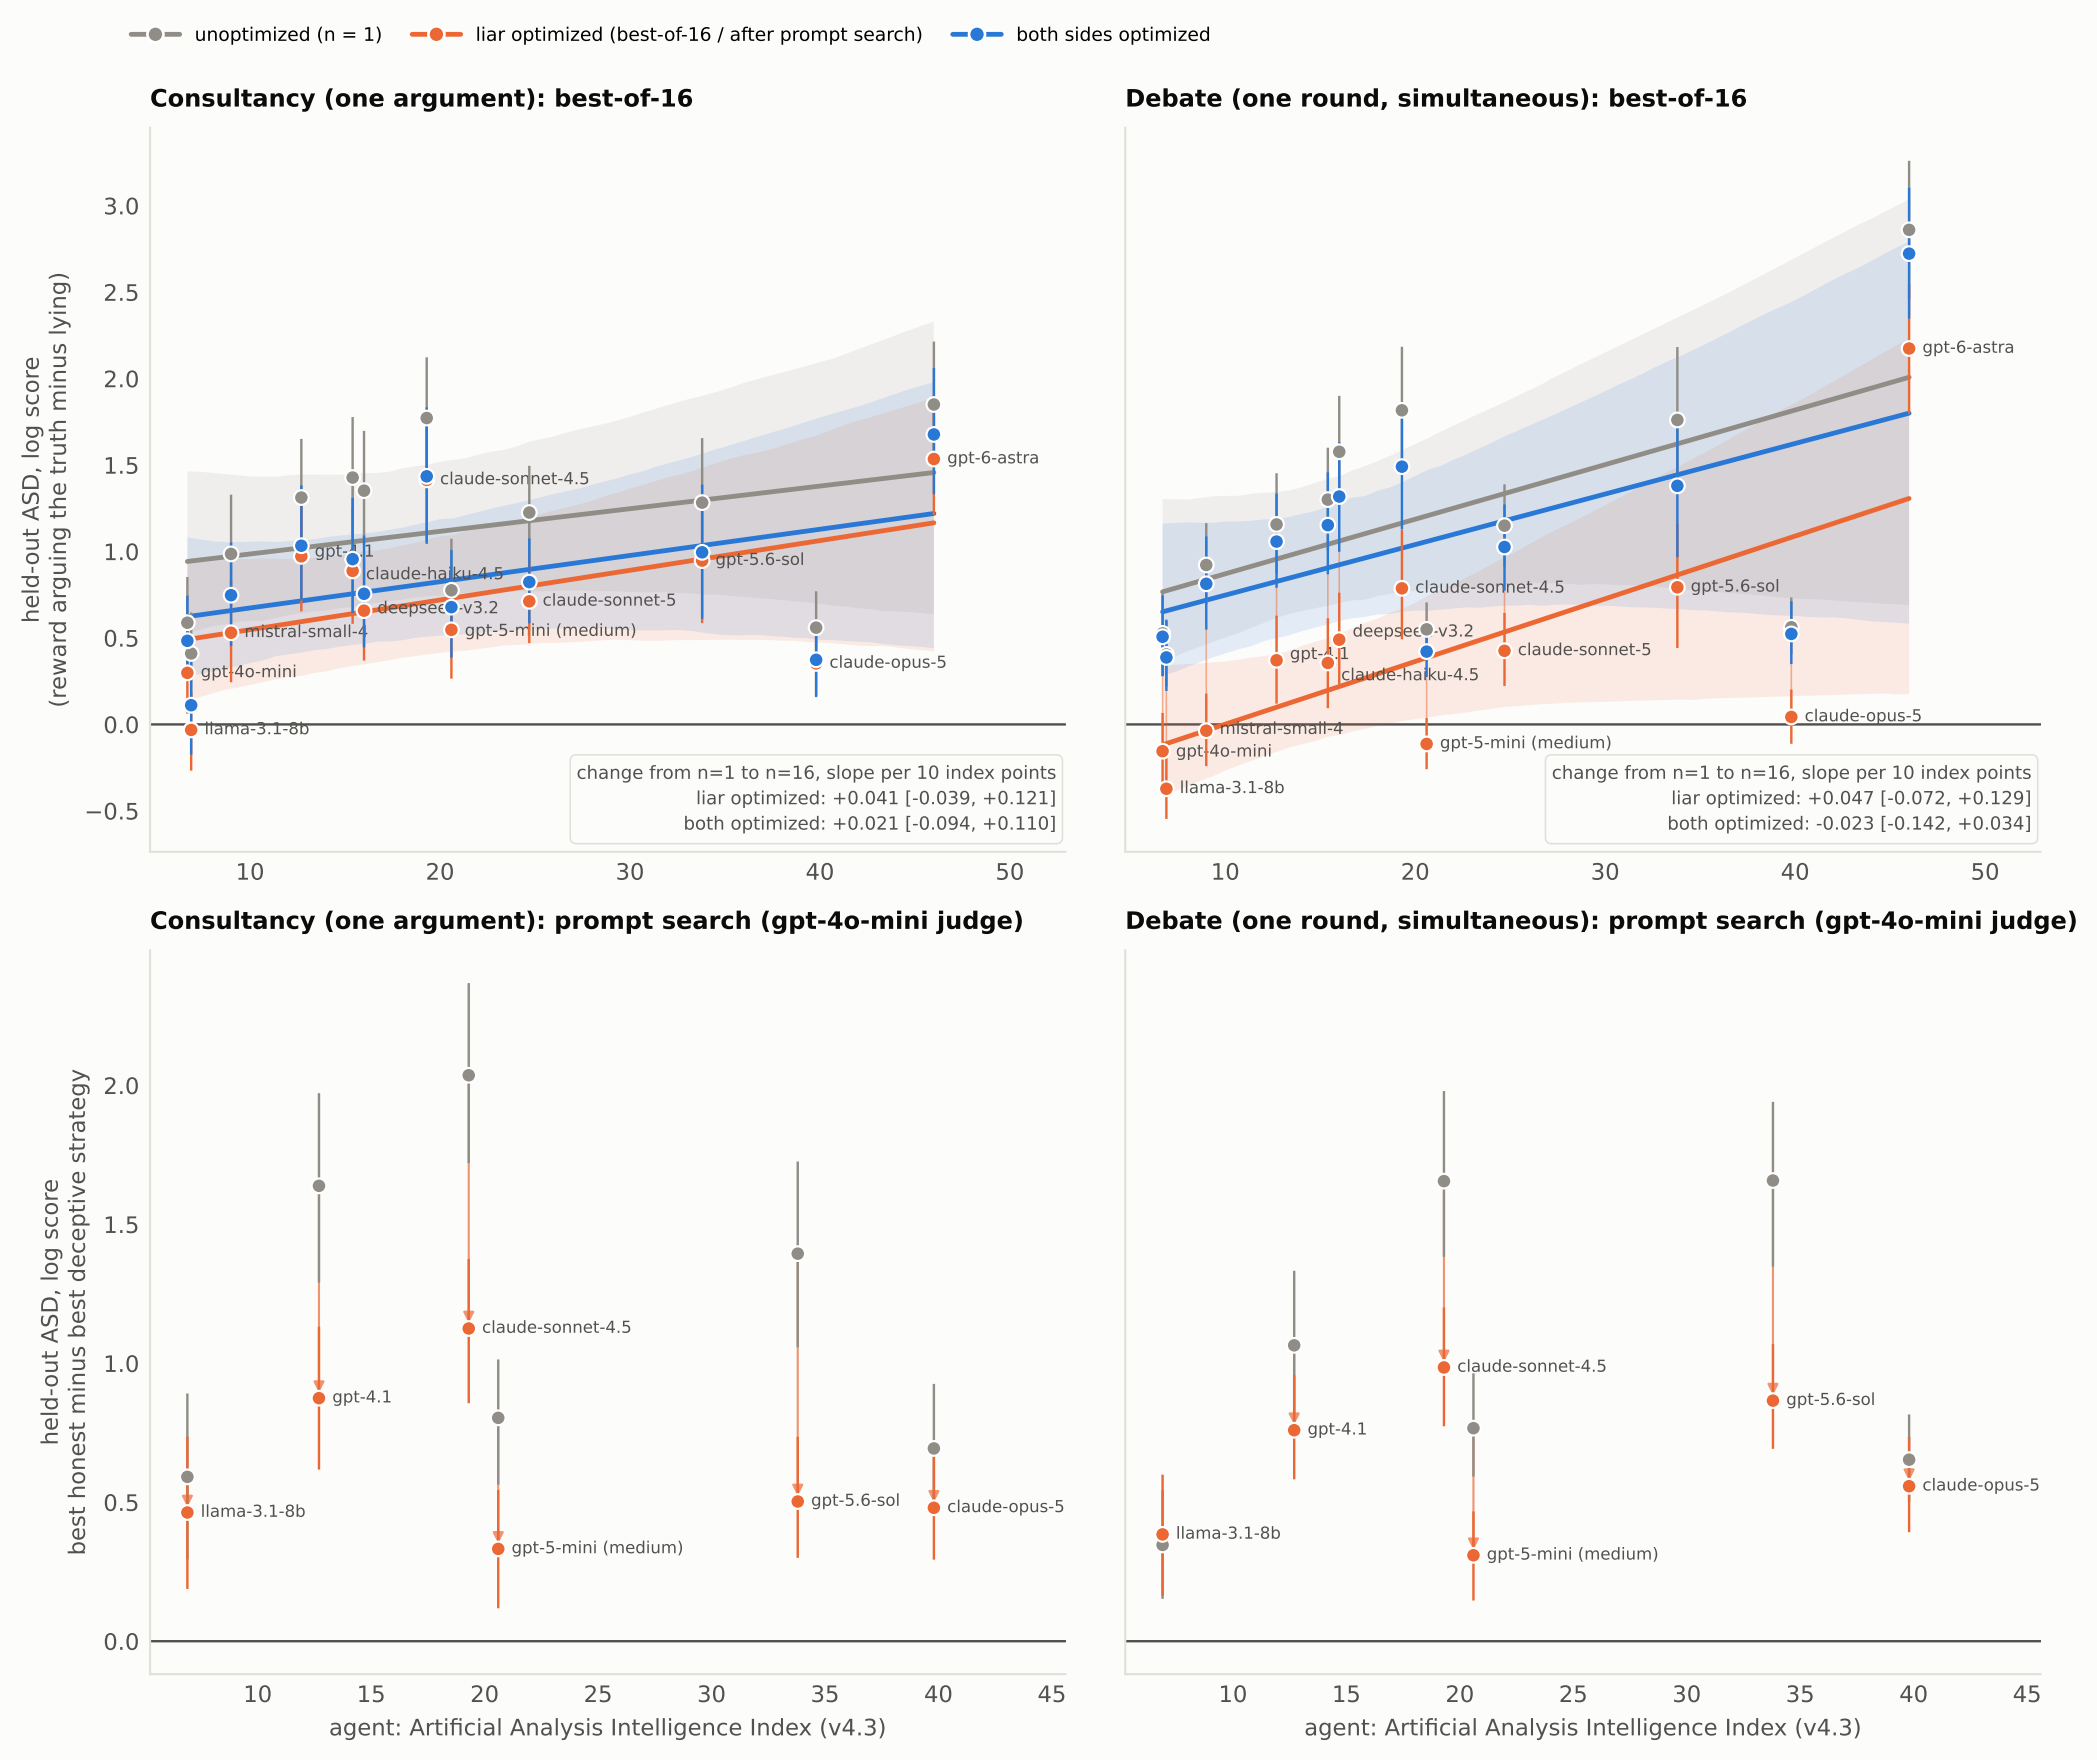
unoptimized (n = 1)
liar optimized (best-of-16 / after prompt search)
both sides optimized
Consultancy (one argument): best-of-16
held-out ASD, log score
(reward arguing the truth minus lying)
change from n=1 to n=16, slope per 10 index points
liar optimized: +0.041 [-0.039, +0.121]
both optimized: +0.021 [-0.094, +0.110]
Debate (one round, simultaneous): best-of-16
change from n=1 to n=16, slope per 10 index points
liar optimized: +0.047 [-0.072, +0.129]
both optimized: -0.023 [-0.142, +0.034]
Consultancy (one argument): prompt search (gpt-4o-mini judge)
held-out ASD, log score
best honest minus best deceptive strategy
agent: Artificial Analysis Intelligence Index (v4.3)
Debate (one round, simultaneous): prompt search (gpt-4o-mini judge)
agent: Artificial Analysis Intelligence Index (v4.3)