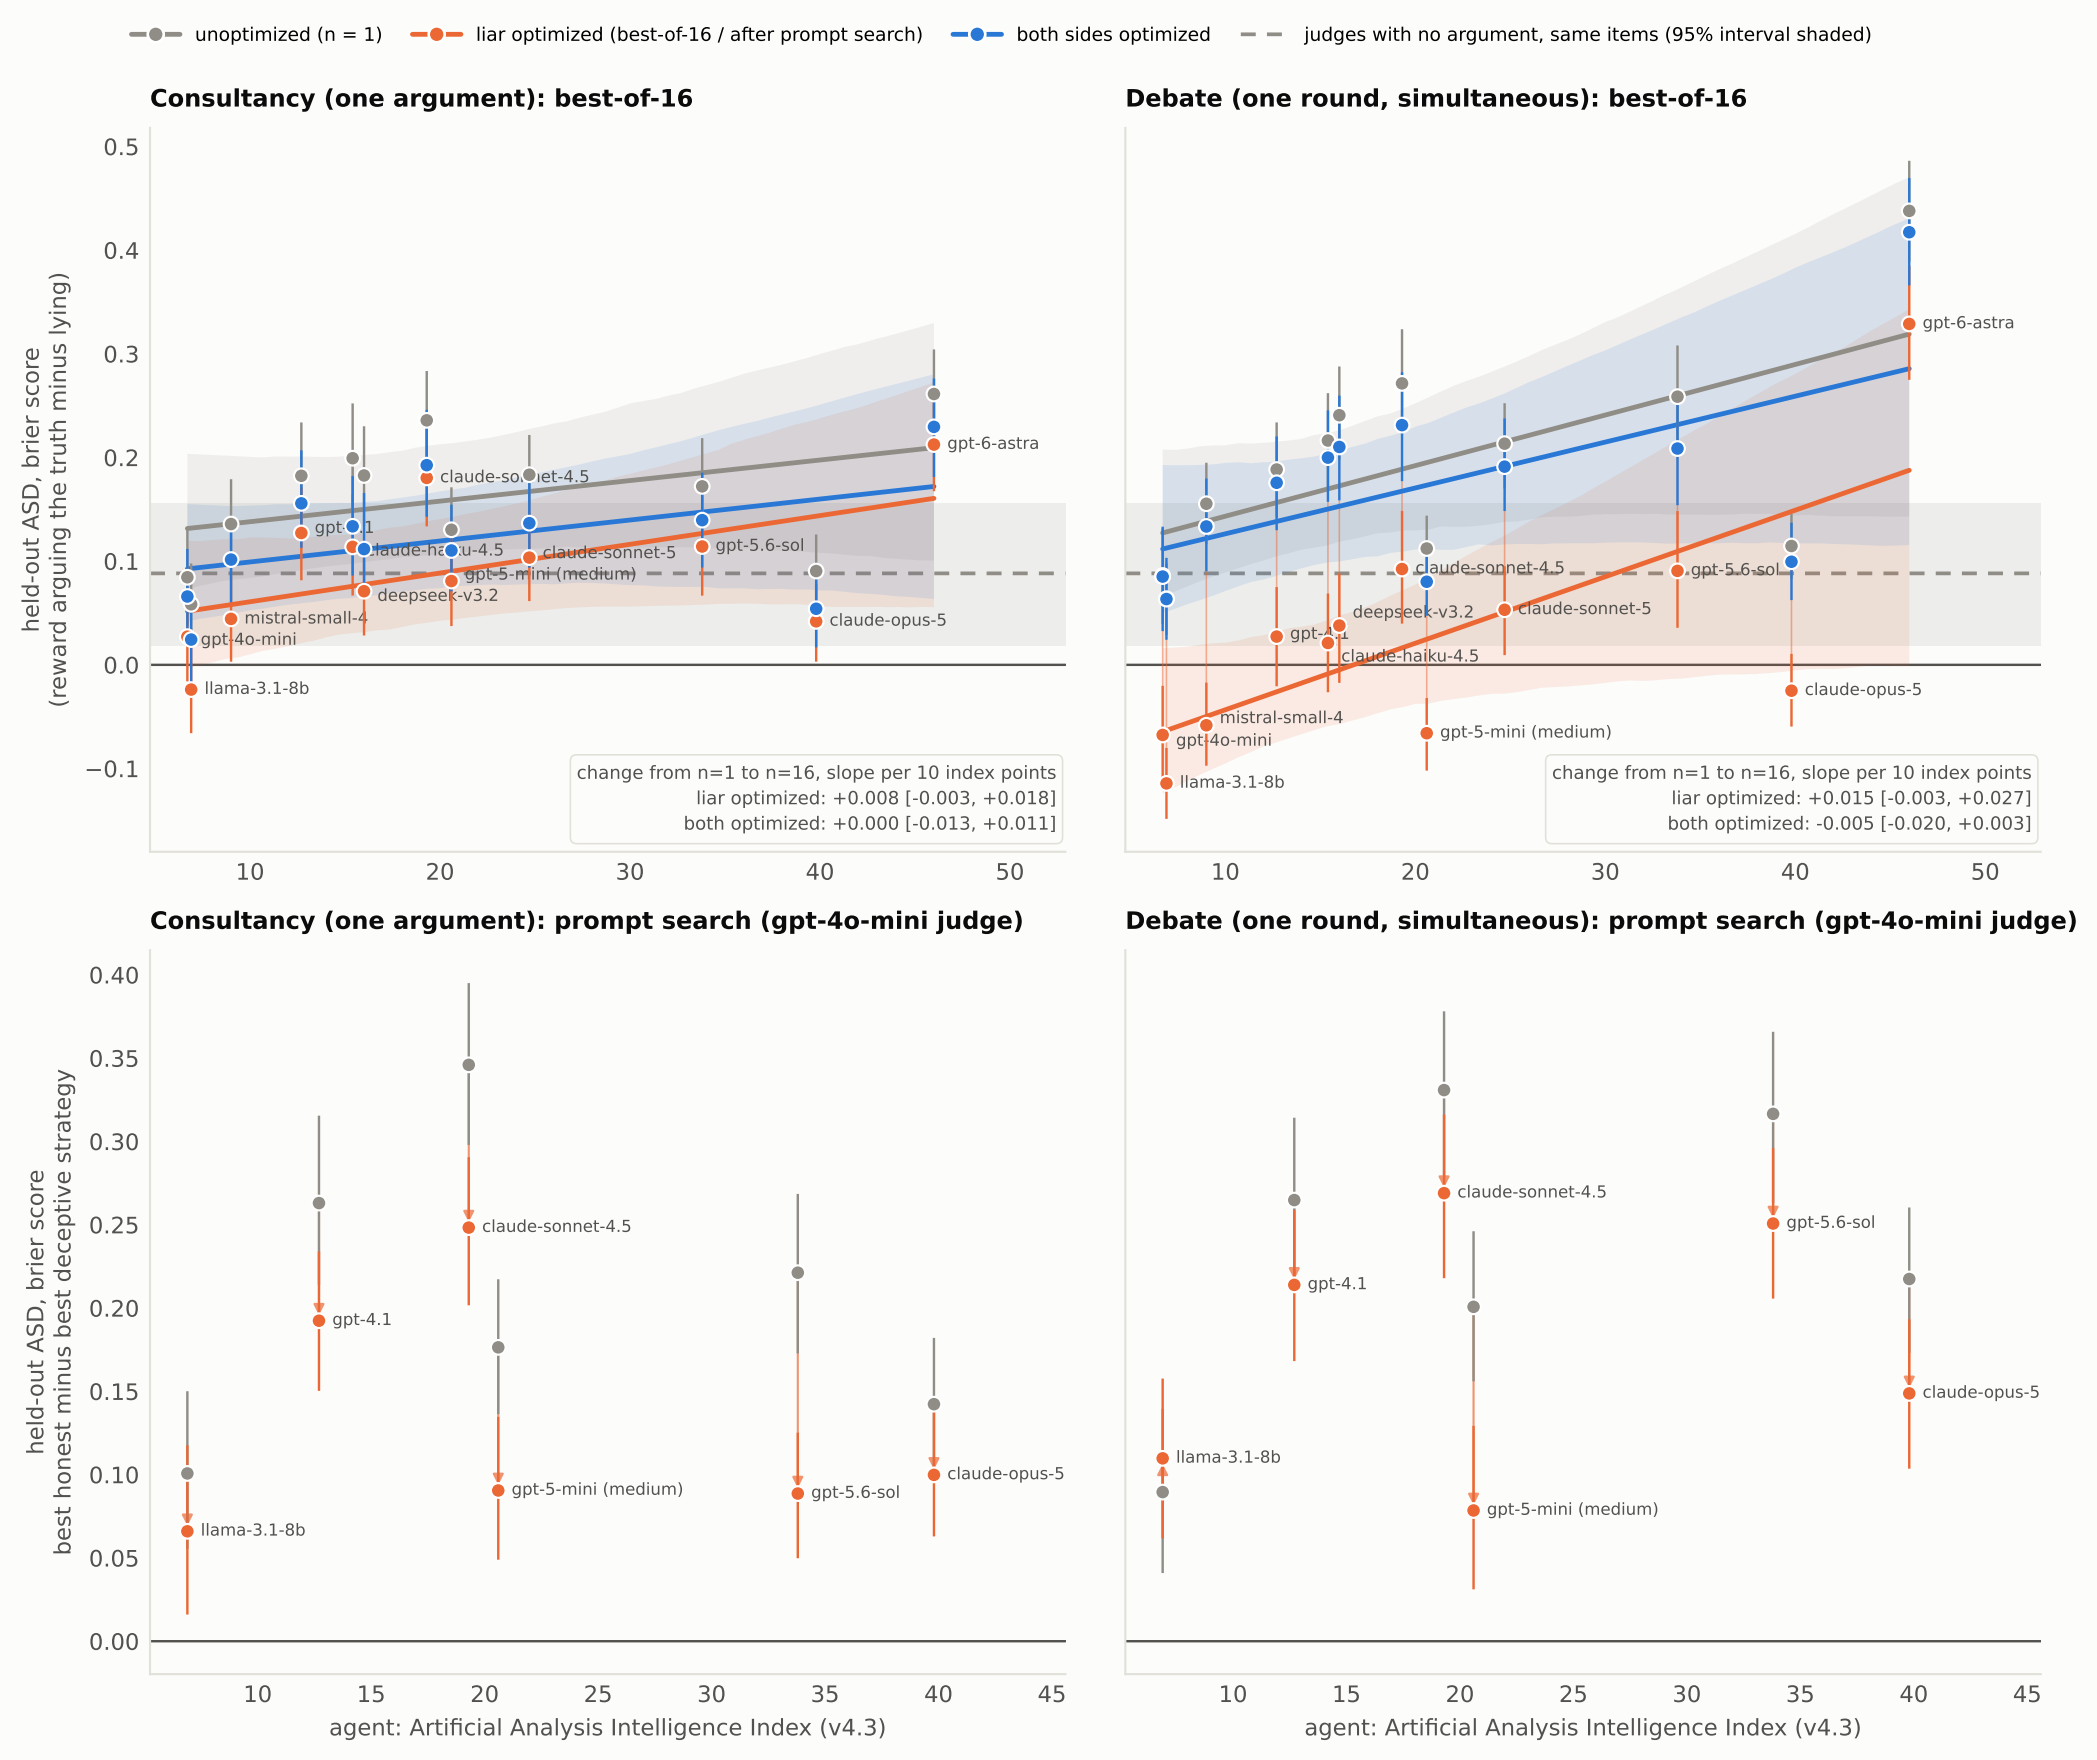
unoptimized (n = 1)
liar optimized (best-of-16 / after prompt search)
both sides optimized
judges with no argument, same items (95% interval shaded)
Consultancy (one argument): best-of-16
Debate (one round, simultaneous): best-of-16
held-out ASD, brier score
(reward arguing the truth minus lying)
gpt-6-astra
claude-sonnet-4.5
gpt-5.6-sol
gpt-5-mini (medium)
deepseek-v3.2
mistral-small-4
gpt-4o-mini
claude-opus-5
llama-3.1-8b
claude-haiku-4.5
change from n=1 to n=16, slope per 10 index points
liar optimized: +0.008 [-0.003, +0.018]
both optimized: +0.000 [-0.013, +0.011]
liar optimized: +0.015 [-0.003, +0.027]
both optimized: -0.005 [-0.020, +0.003]
Consultancy (one argument): prompt search (gpt-4o-mini judge)
Debate (one round, simultaneous): prompt search (gpt-4o-mini judge)
held-out ASD, brier score
best honest minus best deceptive strategy
gpt-4.1
agent: Artificial Analysis Intelligence Index (v4.3)