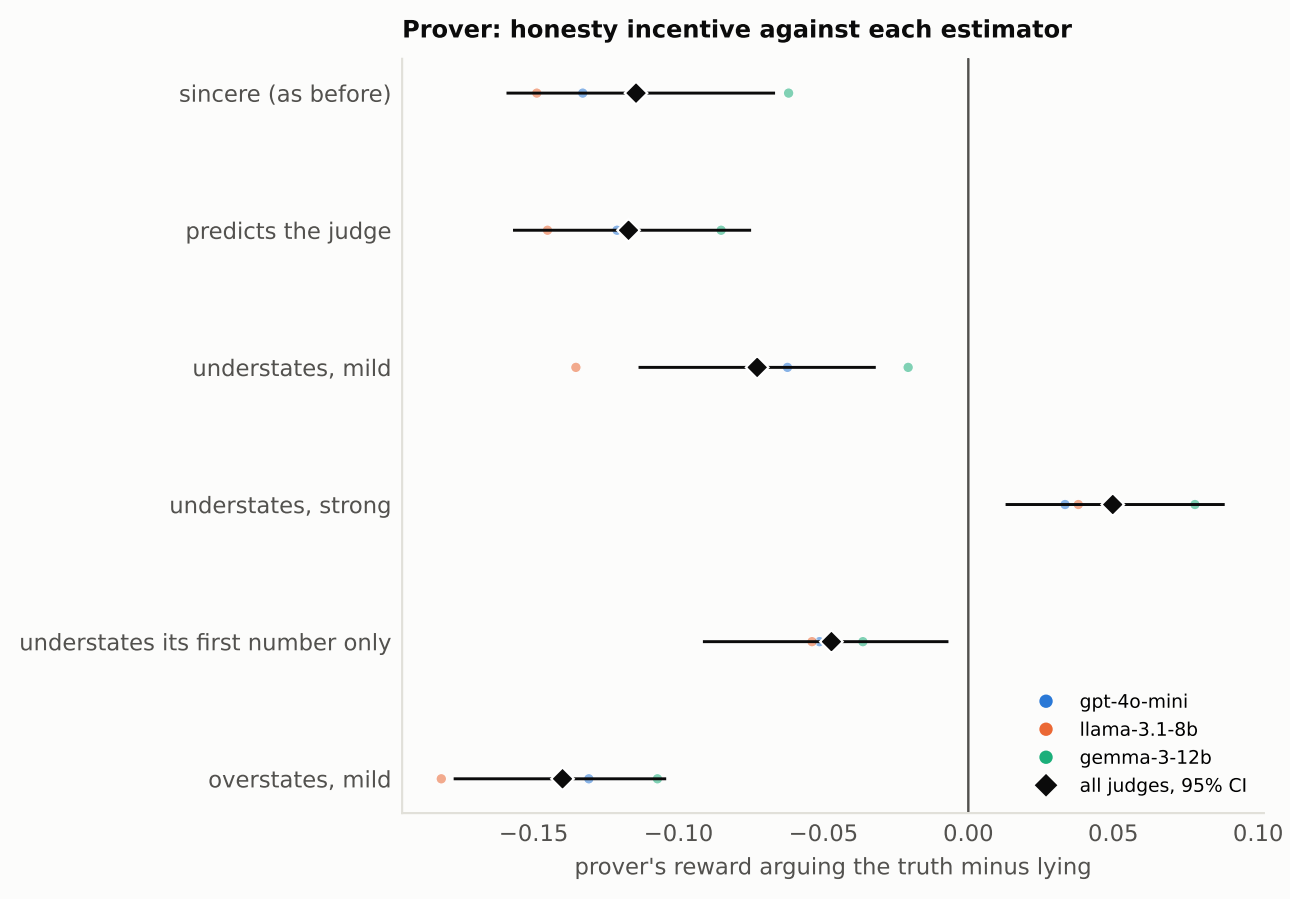
Prover: honesty incentive against each estimator
sincere (as before)
predicts the judge
understates, mild
understates, strong
understates its first number only
overstates, mild
−0.15
−0.10
−0.05
0.00
0.05
0.10
prover's reward arguing the truth minus lying
gpt-4o-mini
llama-3.1-8b
gemma-3-12b
all judges, 95% CI

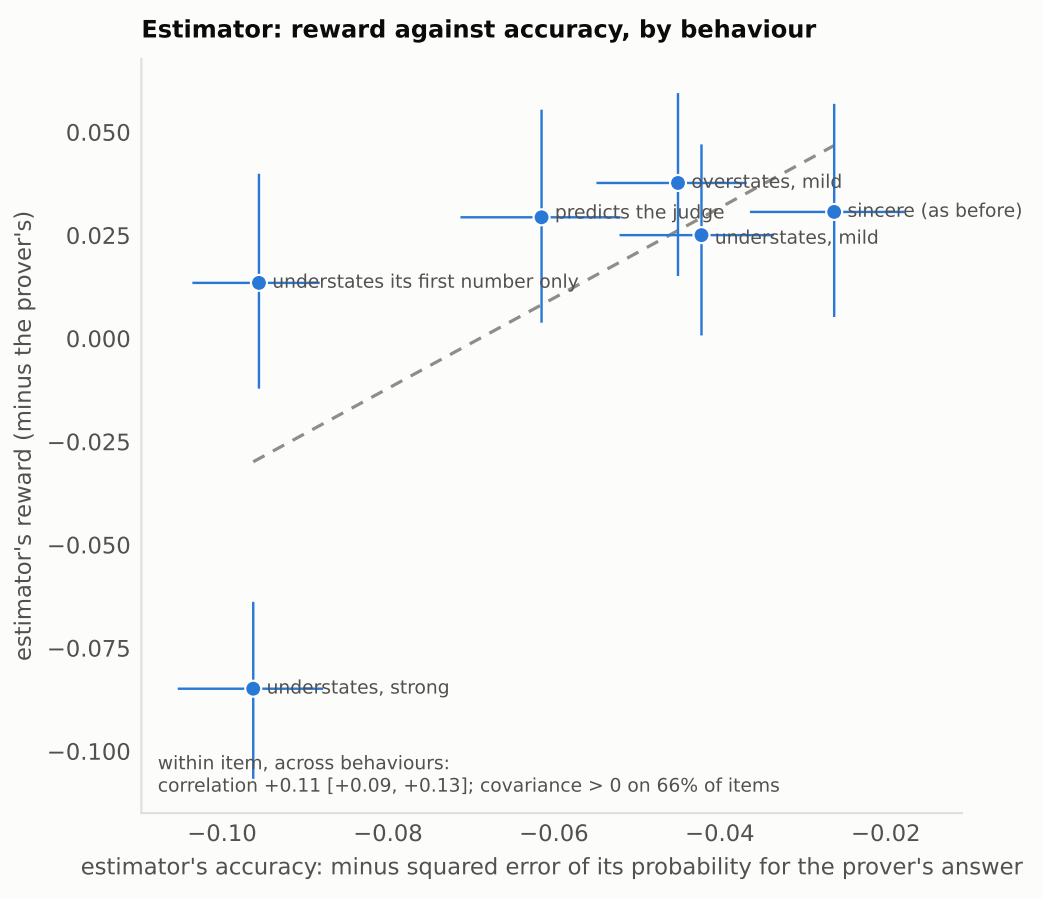
Estimator: reward against accuracy, by behaviour
estimator's reward (minus the prover's)
0.050
0.025
0.000
−0.025
−0.050
−0.075
−0.100
overstates, mild
predicts the judge
sincere (as before)
understates, mild
understates its first number only
understates, strong
within item, across behaviours:
correlation +0.11 [+0.09, +0.13]; covariance > 0 on 66% of items
−0.10
−0.08
−0.06
−0.04
−0.02
estimator's accuracy: minus squared error of its probability for the prover's answer